

Adaptive Learning in Motion: Harnessing Cloud-Based AI for Simulating Smart Navigation

City University of Seattle
School of Technology & Computing
TEAM A

Verónica Elze
ElzeVeronica@CityUniversity.edu
Master of Artificial Intelligence

Vidhyalakshmi Amarnath
AmarnathVidhyalaksh@CityUniversity.edu
Master of Artificial Intelligence

Honorine Ndom Ndzah
NdomndzahHonorine@CityUniversity.edu
Master of Computer Science

Jayasudha Kalamegam
KalamegamJayasudha@CityUniversity.edu
Master of Computer Science

Abstract

This project explores the application of reinforcement learning (RL) techniques and cloud-based tools to train an agent for navigating grid-based environments. Initially conceptualized as a physical “Smart Soccer Ball,” the project evolved into a simulation-based approach using OpenAI Gym for environment generation and AWS for scalable RL training. The Proximal Policy Optimization (PPO) algorithm, implemented through Stable Baselines3, was selected for its reliability in balancing exploration and exploitation during training. The agent’s learning process was conducted in a locally customized MiniGrid environment, where observations were preprocessed into simplified formats for efficient training. Reward shaping guided the agent's behavior by reinforcing goal-reaching actions and penalizing collisions or delays. Key metrics, such as cumulative rewards, navigation time, and success rate, were tracked using TensorBoard to monitor performance. Results demonstrated the agent’s capability to reduce navigation times from an initial 106 seconds to just 2 seconds after iterative training. AWS services such as SageMaker, S3, and CloudWatch enabled scalable data storage, secure access, and performance monitoring. This project highlights the potential of combining RL algorithms with modular, cloud-based systems for navigation tasks and lays a foundation for transitioning to 3D environments and exploring multi-agent systems.

Keywords: reinforcement learning, Proximal Policy Optimization (PPO), AWS SageMaker, OpenAI Gym, virtual agent, autonomous navigation, simulation environments.

1. INTRODUCTION

Problem to Solve

One of the key challenges we face is navigating dynamic environments with an autonomous system. Although autonomous navigation has been explored in various fields, it remains complex when environments constantly change. Reinforcement learning (RL) models, though powerful, often struggle to generalize across varying conditions and adjust to real-time changes (Kober et al., 2013). Traditional RL approaches, which rely heavily on trial and error, can be inefficient in unpredictable environments.

Our project addresses this challenge by training a virtual agent in simulated environments to learn adaptive navigation. Using Proximal Policy Optimization (PPO), a popular RL technique, the agent improves its ability to handle static obstacles in grid-based mazes, laying the groundwork for more complex environments in future phases.

Motivation

One of the key challenges in autonomous navigation is adapting to dynamic environments where conditions constantly change. Reinforcement learning (RL) offers a promising solution, but traditional approaches struggle to generalize across varying scenarios and adjust in real time (Kober et al., 2013). Model-based RL (MBRL) addresses this by enabling systems to anticipate changes, improving adaptability and efficiency (Polydoros & Nalpantidis, 2017). Teaching an autonomous agent to navigate dynamic mazes using RL techniques provides valuable insights into low-cost, scalable AI solutions with real-world applications, including robotics and smart sports technology (Zhang, 2022).

Usefulness/Beneficiaries

This project highlights the potential for AI-driven systems to tackle real-world challenges through simulation-based training. Industries like smart sports technology are particularly interested in low-cost, autonomous solutions that can enhance performance and adaptability (Zhang, 2022).

By successfully training a virtual agent in a simulated environment, we provide insights into how reinforcement learning techniques can be applied to create intelligent, adaptive systems.

From an educational perspective, this project offered hands-on experience with RL methods, AWS SageMaker, and simulation tools like OpenAI

Gym. These skills are increasingly valuable as AI and robotics technologies advance. Beyond this, the project demonstrates how foundational AI techniques can pave the way for real-world applications in navigation, robotics, and performance optimization.

2. LITERATURE REVIEW

Autonomous systems capable of adapting to dynamic environments are increasingly relevant across various domains, including robotics, artificial intelligence (AI), and sports technology. This project aims to train a virtual smart agent to navigate grid-based mazes and static obstacles using RL techniques. This literature review explores foundational research on related methodologies and highlights how these insights shaped the development of this project.

RL in Robotics

RL has emerged as a transformative approach for robotics, allowing systems to adapt dynamically to new scenarios. Kober et al. (2013) provide a comprehensive survey of RL applications in robotics, highlighting its flexibility compared to traditional rule-based systems. Building on this, Mnih et al. (2015) demonstrated how deep reinforcement learning can achieve human-level control, showcasing RL's effectiveness in solving complex decision-making tasks. These foundational insights support the development of RL models for tasks like efficient maze navigation, as applied in this project.

The project specifically benefits from PPO, a model-free RL algorithm. While model-based RL (MBRL) can predict and respond to environmental changes (Polydoros & Nalpantidis, 2017), PPO strikes a balance between performance and computational efficiency, making it well-suited for the grid-based mazes and static obstacles tackled in this project. PPO's adaptability to varying conditions supports the goal of creating a scalable training framework for more complex applications in future phases.

Cloud Computing and AI Scalability

The integration of cloud-based platforms like AWS SageMaker is a pivotal aspect of this project. SageMaker provides a scalable environment for developing, training, and deploying RL models, offering computational resources that can significantly reduce the time and cost of experimentation (Amazon Web Services, 2021). This capability contrasts with earlier robotics

projects that relied heavily on on-premises hardware, which limited scalability.

Tripuraneni and Song (2019) emphasize how cloud infrastructure accelerates AI model development by enabling distributed computing and iterative tuning. SageMaker's capabilities allowed for the rapid training of the PPO algorithm, optimizing navigation strategies through simulation environments built in OpenAI Gym. These insights reinforce the importance of scalability, especially when experimenting with multiple RL approaches or extending to more complex simulations.

Autonomous Navigation and Sensor Integration

Simulated environments form the backbone of many autonomous robotics projects. OpenAI Gym provides a framework for developing RL algorithms, allowing agents to interact with grid-based mazes and progressively improve their performance. Unlike projects that integrate physical sensors, this project focuses on the foundational step of training a virtual agent to adapt to static obstacles efficiently.

Research by Malekzadeh et al. (2018) highlights the importance of sensor-driven decision-making in robotics. While physical sensors were not implemented in this phase of the project, the adaptability demonstrated by the virtual agent aligns with findings on improving obstacle detection through iterative training. Future iterations of the project may incorporate advanced technologies like LIDAR, as Malla and Dholakiya (2022) suggest, to enhance spatial awareness and obstacle navigation in real-world scenarios.

Comparative Analysis with Similar Projects

Zhang (2022) explores AI applications in soccer training, emphasizing the potential of dynamic navigation systems. While Zhang used genetic algorithms to optimize movement paths, this project applies RL techniques that offer greater flexibility. Unlike genetic algorithms, which evolve solutions over multiple iterations, RL models like PPO learn in real-time, making them better suited for dynamic obstacle scenarios.

Shah et al. (2020) introduced the AirSim framework for training autonomous vehicles using high-fidelity simulations. Similarly, this project leverages OpenAI Gym for grid-based simulations, emphasizing lightweight, cost-effective approaches compared to resource-intensive solutions like AirSim. This focus ensures

scalability for future advancements in adaptive robotics.

Integration and Contribution

The project synthesizes advances in RL and cloud-based AI to train a virtual agent capable of navigating static environments. By leveraging PPO, AWS SageMaker, and OpenAI Gym, the project demonstrates how foundational RL techniques can address challenges in autonomous navigation. Compared to related work, the project emphasizes cost-effectiveness and scalability, providing a platform for future research in dynamic environments and physical robotics applications.

This literature review highlights the foundational research and methodologies that guide the project's development, positioning it as a meaningful step in exploring adaptive AI systems within simulated environments.

3. METHODOLOGY

This project focuses on training a virtual agent to navigate grid-based mazes using RL. By prioritizing local development and validation, we ensured a robust foundation for the agent's training and evaluation in simulated environments before scaling to cloud-based platforms for future iterations.

Environment Setup and Customization

Using the MiniGrid library, we developed a grid-based maze environment tailored for RL tasks. A custom observation wrapper was implemented to preprocess and flatten the multi-dimensional observations into a single vector, streamlining compatibility with RL algorithms. Discrete actions, including movements in cardinal directions and interactions with the environment, were defined to enable efficient navigation by the agent. All customization and validation of the environment have been carried out in a local environment, ensuring that the setup is robust before integration with external systems.

Agent Training

The PPO algorithm was used to train the agent, leveraging the Stable Baselines3 framework. PPO is widely recognized for its balance between sample efficiency and ease of implementation (Schulman et al., 2017), making it a suitable choice for our grid-based maze environment. Training was executed locally, utilizing the DummyVecEnv wrapper to vectorize the environment and enable efficient agent-

environment interactions. The PPO algorithm was chosen for its simplicity and effectiveness, while alternative algorithms like Soft Actor-Critic (SAC) (Haarnoja et al., 2018) could be explored in future iterations for improved stability and entropy maximization.

Training progress was monitored with TensorBoard to track real-time metrics such as rewards, policy loss, and training duration. The agent's task involved navigating the grid-based environment to reach the green goal square while avoiding obstacles. *Appendix E* illustrates the agent's behavior during training, showing the agent (red triangle) progressing toward the goal (green square) in the customized MiniGrid environment.

By logging metrics like reward progression and training loss, we validated the agent's learning progress. Through reward shaping and iterative fine-tuning, the agent's navigation time improved significantly, decreasing from an initial 106 seconds to just 2 seconds by the end of training. Detailed logs and outputs are available in *Appendix B* and *Appendix C*, showcasing the incremental performance improvements.

Evaluation and Validation

Our local setup supported detailed evaluation and debugging of both the environment and agent behavior. Observations were validated to ensure accurate preprocessing, and actions were tested to confirm expected outcomes. Debugging efforts focused on resolving compatibility issues, such as observation flattening errors and reward signal inconsistencies. The agent's progress was measured by its ability to optimize navigation time, decreasing from approximately 106 seconds to 2 seconds by the end of training.

Specific local environment outputs, including action logs, episode completions, and rewards, are provided in *Appendix C*. These results demonstrate the agent's learning behavior during local development.

Following successful local validation, the model was transitioned to AWS SageMaker for further training and evaluation. SageMaker's robust infrastructure enabled monitoring of key metrics such as entropy loss, explained variance, and reward progression while securely storing model checkpoints in Amazon S3. Outputs from the

SageMaker instance, including training logs and final model results, are detailed in *Appendix D*.

Tools

The success of this project relied on a suite of software tools for RL development, training, and monitoring.

Software Tools

- Python Programming Language: Provides the foundation for RL model development, cloud integration, and environment customization.
- MiniGrid Library: A minimalistic, grid-based environment used for creating and customizing maze scenarios for training our RL agent.
- Stable Baselines3 Framework: Facilitates the implementation of the PPO algorithm, chosen for its efficiency and reliability in RL tasks.
- OpenAI Gym: Allows preliminary testing in a simulated environment to refine RL strategies before deployment.

Cloud Tools (to be implemented)

- AWS SageMaker: Used to train and fine-tune the RL model. Its scalable infrastructure allowed efficient optimization of the PPO algorithm, reducing the time and cost of local experimentation.
- Amazon S3: Acted as the centralized storage hub for simulation data and model checkpoints, ensuring secure and accessible data management throughout the training process.
- AWS CloudWatch: Provided real-time monitoring of system resources and application logs during training, enabling proactive issue identification and system performance optimization.
- AWS Identity and Access Management (IAM): Ensured secure access to cloud resources, maintaining data privacy and workflow integrity during model training and experimentation.

Development and Monitoring Tool

- TensorBoard: Visualized training logs in real-time, including metrics like rewards, policy loss, and episode durations. This helped track the agent's progress, fine-tune hyperparameters, and validate the effectiveness of the PPO algorithm.

Technical Architecture



Figure 1 3D perspective technical architecture diagram for RL project by DALL·E (OpenAI, 2024).

Figure 1 illustrates the flow of data through the system: OpenAI Gym generates the simulation environment, AWS IAM secures access, S3 stores the training data, SageMaker trains the RL agent, and CloudWatch monitors performance metrics. This architecture leverages cloud scalability to optimize the RL training process efficiently. For an enlarged version of the 3D image showcasing the architecture in greater detail, see Appendix A.

Our project leverages a combination of OpenAI Gym and AWS to create a dynamic, scalable, and secure training environment for our smart agent. Here's how it works:

- **Simulation Environment:** OpenAI Gym serves as our digital soccer field, providing a simulated environment where the agent can learn and practice. This is where the agent interacts with the ball and its surroundings.
- **Secure Data Flow:** AWS Identity and Access Management (IAM) acts as a security layer, controlling access to the system and ensuring that only authorized users and services can interact with the training process.
- **Data Storage:** Amazon S3, our secure and scalable storage service, stores all the important data generated during training. This includes simulation states, actions taken by the agent, and rewards received.
- **Agent Training:** Amazon SageMaker, our machine learning platform, takes the data from S3 and uses it to train the smart agent. SageMaker enables the agent to learn strategies and make intelligent decisions,

transforming it from a novice to an expert player.

- **Monitoring and Logging:** Amazon CloudWatch acts as a monitoring and logging service, keeping track of performance metrics, recording results, and alerting us to any issues that may arise during training.

This combination of OpenAI Gym and AWS services establishes a robust architecture for training and evaluating our RL agent. The architecture's scalability and integration of cloud resources streamline the development process, ensuring an efficient and optimized workflow.

With the technical foundation established and training processes implemented, the next step was to evaluate the model's performance using key metrics. The following section, Metrics and Visualizations, highlights the agent's learning progress, behavior, and efficiency through detailed analysis and visual representations.

4. METRICS AND VISUALIZATIONS

Metrics are quantitative measures used to evaluate the performance and effectiveness of the RL model. These metrics provide insights into the model's learning progress, adaptability, and computational efficiency. To evaluate the performance of the RL model, several key metrics were employed, categorized as follows:

Testing Metrics:

- **Fluctuating Rewards:** The graph highlights inconsistent agent performance, with rewards varying significantly across tests.
- **Highs and Lows:** Some tests achieved near-optimal results, while others showed minimal progress, indicating instability.
- **Stabilization Challenges:** Results suggest the need for further fine-tuning to address the environmental complexity.




Figure 2 Total Reward across multiple tests

Training Metrics:

- Cumulative Rewards: This metric tracks the improvement in the model's performance over time, providing insight into the agent's learning progress.
- Convergence Time: Measures the time taken for the RL model to stabilize and achieve consistent performance.

Simulation Metrics:

- Success Rate: The percentage of successful navigations, indicating the agent's effectiveness in reaching its goals.
- Obstacle Avoidance Accuracy: Evaluates the agent's ability to detect and avoid obstacles during navigation.

Efficiency Metrics:

- Training Time per Epoch: Captures the time required to complete each training iteration, reflecting the computational speed.
- Simulation Runtime: Assesses the overall efficiency of the simulation environment during testing and evaluation.

Robustness Metrics:

- Generalization: Tests the model's adaptability and performance in unseen environments.
- Failure Rate: Tracks the frequency of undesirable outcomes, such as collisions or incomplete tasks.

Visualization

Visualization plays a vital role in interpreting and demonstrating the RL model's performance. During the training, progress, and agent behaviors were visualized. Key visual elements include:

- Graphs show the increase in cumulative rewards over time as the agent learns.
- Visual demonstrations of the agent navigating simulated environments, highlighting its interactions with obstacles and goals.




Figure 3 Mean Reward during Evaluation

The "eval/mean_reward" graph (Figure 3), is a crucial visualization of our agent's performance. It shows how the average reward earned changed over time during the evaluation phase. Upward trend clearly demonstrates the agent's learning progress and its ability to achieve increasingly higher rewards as it gains experience. This indicates the agent is not only successfully completing the task but also becoming more efficient and strategic in its actions.

Appendix F & G provide a more comprehensive view of the model's performance during training. Below are descriptions of some of them that help provide a comprehensive view of the agent's learning process:

- **Increasing Rewards (rollout/ep_reward_mean, eval/mean_reward):** These graphs demonstrate the agent's ability to learn and improve its performance, achieving higher rewards as training progresses.

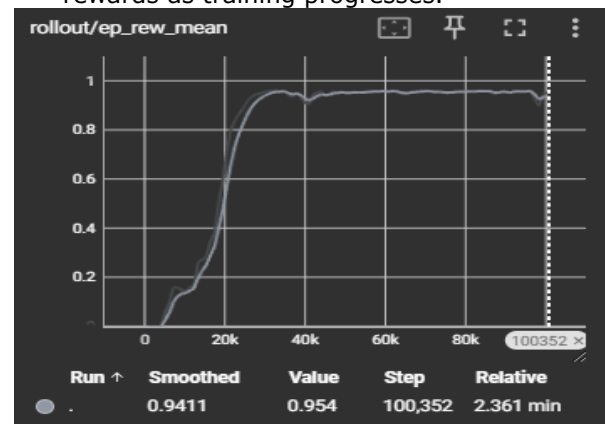


Figure 4 Increasing Rewards during training

- **Decreasing Loss (train/loss):** This graph shows a general downward trend, indicating that the agent's policy and value function are being effectively optimized during training.

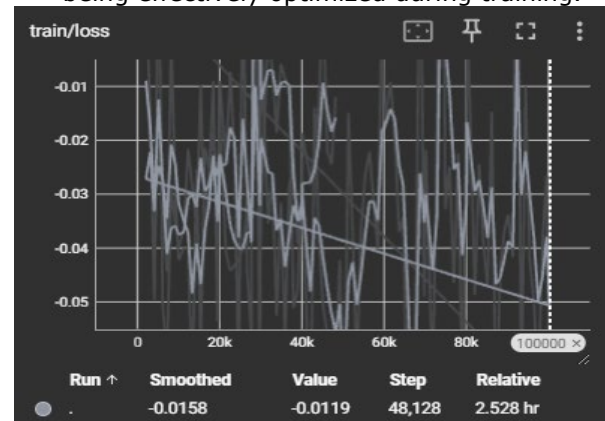


Figure 5 Decreasing Loss during training

- **Stable Learning (train/approx_kl, train/clip_fraction):** These graphs suggest that the training process was stable, with the agent exploring new strategies without making drastic, destabilizing changes to its policy.

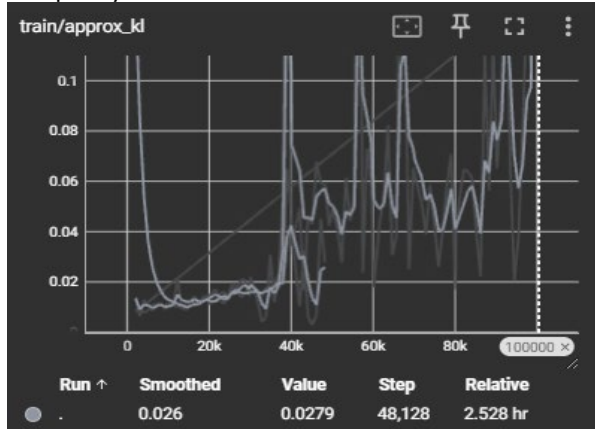


Figure 6 Stable Learning during training

- **Improved Value Estimation (train/explained_variance):** This graph shows that the agent is learning to accurately predict the value of different states and actions, which is crucial for making informed decisions.

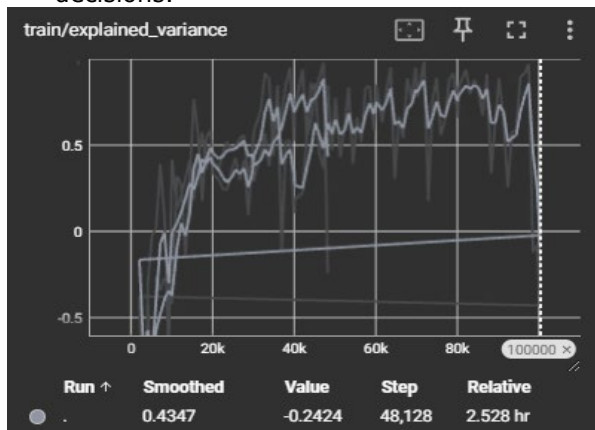


Figure 7 Improved Value Estimation during training

With these metrics and visualizations, we can clearly observe the agent's learning trends, performance improvements, and areas for optimization. The next section, Results, discusses these findings in detail, focusing on how the RL model performed across training and testing phases.

5. RESULTS

Results demonstrated the agent's ability to optimize navigation through iterative training

with PPO. Navigation times decreased from an initial 106 seconds to just 2 seconds, showcasing improved efficiency and learning progress. Metrics such as cumulative rewards and episode success rates confirmed the agent's generalization across varied maze configurations.

6. CONCLUSION

This project successfully demonstrated the potential of reinforcement learning combined with cloud-based tools for solving navigation tasks in dynamic environments. By leveraging PPO and scalable AWS services, the agent achieved efficient learning and navigation performance. Future work can expand these foundations by exploring 3D environments and multi-agent systems.

While this project successfully established a robust foundation, the following section explores opportunities for future enhancements.

7. FUTURE WORK

While our project achieved its primary objectives, there remains significant opportunity to expand and refine the system. The next phase could focus on building upon our progress to address more complex challenges and enhance overall functionality.

Focusing on a transition from a 2D simulation to a 3D environment would be a natural progression, adding depth and complexity to the agent's learning process. By incorporating 3D simulations, the agent could be trained to handle elevation changes, varied terrain, and dynamic lighting conditions, making its navigation capabilities more realistic and robust. Advanced technologies such as SimSpace Weaver could facilitate these 3D simulations, enabling real-time interaction with dynamic objects and more sophisticated scenarios. Future iterations could also explore multi-agent RL to enable collaborative behaviors or competitive dynamics, further expanding the system's practical applications in real-world scenarios such as robotic soccer or autonomous exploration missions.

8. ACKNOWLEDGEMENTS

We would like to acknowledge the assistance of OpenAI's ChatGPT tool in improving the verbiage, flow, and transitions of this paper as well as offering content such as the 3D DALL-E image.

9. REFERENCES

The following references include a combination of sources recommended by ChatGPT, with over half published or updated within the last five years. All citations adhere to APA (American Psychological Association) guidelines.

Scholarly Sources

These sources provide foundational research and insights relevant to the project.

- Amazon Web Services. (2021). *AWS SageMaker: Developer Guide*. Retrieved from <https://docs.aws.amazon.com/sagemaker/latest/dg/whatis.html>
- Haarnoja, T., Zhou, A., Abbeel, P., & Levine, S. (2018). Soft actor-critic: Off-policy maximum entropy deep reinforcement learning with a stochastic actor. *Proceedings of the 35th International Conference on Machine Learning (ICML)*, 2961–2970. <https://arxiv.org/abs/1801.01290>
- Kober, J., Bagnell, J. A., & Peters, J. (2013). Reinforcement learning in robotics: A survey. *The International Journal of Robotics Research*, 32(11), 1238–1274. <https://doi.org/10.1177/0278364913495721>
- Malla, A., & Dholakiya, M. (2022). Autonomous navigation in unstructured environments using LIDAR-based reinforcement learning. *International Journal of Advanced Robotic Systems*, 19(1), 1–12. <https://doi.org/10.1177/17298814221075422>
- Mnih, V., Kavukcuoglu, K., Silver, D., Rusu, A. A., Veness, J., Bellemare, M. G., ... & Hassabis, D. (2015). Human-level control through deep reinforcement learning. *Nature*, 518(7540), 529–533. <https://doi.org/10.1038/nature14236>
- OpenAI. (2024). *3D perspective technical architecture diagram for reinforcement learning project* [3D image]. DALL·E.
- Polydoros, A. S., & Nalpantidis, L. (2017). Survey of model-based reinforcement learning: Applications on robotics. *Journal of Intelligent & Robotic Systems*, 86(2), 153–173. <https://doi.org/10.1007/s10846-017-0570-0>
- Schulman, J., Wolski, F., Dhariwal, P., Radford, A., & Klimov, O. (2017). Proximal policy optimization algorithms. *arXiv Preprint*. <https://arxiv.org/abs/1707.06347>
- Shah, S., Dey, D., Lovett, C., & Kapoor, A. (2020). AirSim: High-fidelity visual and physical simulation for autonomous vehicles. *Conference on Autonomous Agents and Multiagent Systems (AAMAS)*. Retrieved from <https://github.com/microsoft/AirSim>
- Tripuraneni, V., & Song, M. (2019). Cloud infrastructure for scalable machine learning: A comprehensive study. *Journal of Cloud Computing*, 8(1), 1–18. <https://doi.org/10.1186/s13677-019-0124-9>
- Zhang, Y. (2022). Genetic algorithms in sports: Enhancing soccer strategy through artificial intelligence. *Computers in Sport Science*, 11(3), 215–231. <https://doi.org/10.1016/j.coss.2022.07.003>

APPENDIX A

Technical Architecture

This image, created by DALL-E, is an enlarged version of the Technical Architecture diagram from the paper, illustrating the flow of data and interaction between OpenAI Gym and AWS services in greater visual detail.



APPENDIX B

Local Environment Initialization, Agent Interaction, and Episode Completion Output

This appendix includes visualizations from TensorBoard logs, showcasing key training metrics such as reward progression and policy loss over time.

```
C:\Users\MissV\OneDrive\Documents\Education\CityU\2024FallQ4\AI620\Code> python environment.py
Custom observation shape: (1, 193) Observation space: Box(-inf, inf, (193,), float32)
Action space: Discrete(7) C:\Users\MissV\OneDrive\Documents\Education\CityU\2024FallQ4\AI620\Code\venv\Lib\site-
packages\stable_baselines3\common\vec_env\base_vec_env.py:243: UserWarning: You tried to call render() but no render_mode
was passed to the env constructor. warnings.warn("You tried to call render() but no render_mode was passed to the env
constructor.")
Action taken: [3] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [2] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [4] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [3] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [4] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [6] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [2] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [5] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [0] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [2] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [5] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [2] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [3] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [4] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [3] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [6] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [4] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [2] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [5] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [5] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [6] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [6] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [5] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [6] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [3] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [3] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [2] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [2] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [6] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [0] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [2] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [3] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [2] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [4] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [5] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [0] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [3] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [3] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [5] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [4] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [6] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [4] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [2] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [6] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [0] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [6] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [6] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [2] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [4] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
Action taken: [1] Reward: [0.], Done: [False], Info: [{'TimeLimit.truncated': False}]
```

[illegible]

[illegible]

APPENDIX C

Agent Training

This appendix provides a still image of the environment during training in human render mode, illustrating the agent's perspective within the grid-world maze.

rollout/		
ep_len_mean	12.9	
ep_rew_mean	0.955	
time/		
fps	723	
iterations	54	
time_elapsed	76	
total_timesteps	55296	
train/		
approx_kl	0.055399723	
clip_fraction	0.155	
clip_range	0.2	
entropy_loss	-0.167	
explained_variance	0.661	
learning_rate	0.0003	
loss	-0.0264	
n_updates	530	
policy_gradient_loss	-0.0259	
value_loss	0.000623	

rollout/		
ep_len_mean	12.2	
ep_rew_mean	0.957	
time/		
fps	723	
iterations	55	
time_elapsed	77	
total_timesteps	56320	
train/		
approx_kl	0.27952486	
clip_fraction	0.146	
clip_range	0.2	
entropy_loss	-0.129	
explained_variance	0.52	
learning_rate	0.0003	
loss	-0.0221	
n_updates	540	
policy_gradient_loss	-0.0178	
value_loss	0.000592	

rollout/		
ep_len_mean	11.9	
ep_rew_mean	0.958	
time/		
fps	720	
iterations	56	
time_elapsed	79	
total_timesteps	57344	
train/		
approx_kl	0.023819925	
clip_fraction	0.24	
clip_range	0.2	
entropy_loss	-0.217	
explained_variance	0.833	
learning_rate	0.0003	
loss	-0.0315	
n_updates	550	
policy_gradient_loss	-0.0296	
value_loss	0.000192	

rollout/		
ep_len_mean	12.3	
ep_rew_mean	0.957	
time/		

fps	720	
iterations	57	
time_elapsed	81	
total_timesteps	58368	
train/		
approx_kl	0.073045135	
clip_fraction	0.135	
clip_range	0.2	
entropy_loss	-0.154	
explained_variance	0.843	
learning_rate	0.0003	
loss	-0.04	
n_updates	560	
policy_gradient_loss	-0.0231	
value_loss	0.000152	

rollout/		
ep_len_mean	11.6	
ep_rew_mean	0.959	
time/		
fps	720	
iterations	58	
time_elapsed	82	
total_timesteps	59392	
train/		
approx_kl	0.06047496	
clip_fraction	0.109	
clip_range	0.2	
entropy_loss	-0.131	
explained_variance	0.589	
learning_rate	0.0003	
loss	0.00682	
n_updates	570	
policy_gradient_loss	-0.0187	
value_loss	0.000467	

Eval num_timesteps=60000, episode_reward=0.96 +/- 0.00
Episode length: 11.00 +/- 0.00

eval/		
mean_ep_length	11	
mean_reward	0.961	
time/		
total_timesteps	60000	
train/		
approx_kl	0.017942129	
clip_fraction	0.043	
clip_range	0.2	
entropy_loss	-0.094	
explained_variance	0.852	
learning_rate	0.0003	
loss	-0.0165	
n_updates	580	
policy_gradient_loss	-0.0123	
value_loss	0.000154	

rollout/		
ep_len_mean	13	
ep_rew_mean	0.954	
time/		
fps	718	
iterations	59	
time_elapsed	84	
total_timesteps	60416	

rollout/		
ep_len_mean	12	
ep_rew_mean	0.958	
time/		
fps	718	
iterations	60	
time_elapsed	85	
total_timesteps	61440	
train/		
approx_kl	0.0436577	
clip_fraction	0.262	
clip_range	0.2	
entropy_loss	-0.248	
explained_variance	0.632	
learning_rate	0.0003	
loss	-0.0092	
n_updates	590	
policy_gradient_loss	-0.0319	
value_loss	0.000525	

rollout/		
ep_len_mean	11.3	
ep_rew_mean	0.96	
time/		
fps	719	
iterations	61	
time_elapsed	86	
total_timesteps	62464	
train/		
approx_kl	0.054237213	
clip_fraction	0.0599	
clip_range	0.2	
entropy_loss	-0.106	
explained_variance	0.818	
learning_rate	0.0003	
loss	-0.019	
n_updates	600	
policy_gradient_loss	-0.0169	
value_loss	0.000258	

rollout/		
ep_len_mean	14.7	
ep_rew_mean	0.948	
time/		
fps	718	
iterations	62	
time_elapsed	88	
total_timesteps	63488	
train/		
approx_kl	0.08230614	
clip_fraction	0.37	
clip_range	0.2	
entropy_loss	-0.3	
explained_variance	0.951	
learning_rate	0.0003	
loss	-0.0369	
n_updates	610	
policy_gradient_loss	0.023	
value_loss	2.75e-05	

rollout/		
ep_len_mean	15.1	
ep_rew_mean	0.947	
time/		
fps	718	
iterations	63	
time_elapsed	89	
total_timesteps	64512	
train/		
approx_kl	0.03069654	
clip_fraction	0.202	
clip_range	0.2	
entropy_loss	-0.277	

explained_variance	0.306	
learning_rate	0.0003	
loss	-0.0343	
n_updates	620	
policy_gradient_loss	-0.0198	
value_loss	0.000805	

rollout/		
ep_len_mean	15	
ep_rew_mean	0.947	
time/		
fps	719	
iterations	64	
time_elapsed	91	
total_timesteps	65536	
train/		
approx_kl	0.03838364	
clip_fraction	0.217	
clip_range	0.2	
entropy_loss	-0.263	
explained_variance	0.667	
learning_rate	0.0003	
loss	-0.073	
n_updates	630	
policy_gradient_loss	-0.0305	
value_loss	0.000518	

rollout/		
ep_len_mean	12.2	
ep_rew_mean	0.957	
time/		
fps	718	
iterations	65	
time_elapsed	92	
total_timesteps	66560	
train/		
approx_kl	0.19488329	
clip_fraction	0.347	
clip_range	0.2	
entropy_loss	-0.282	
explained_variance	0.54	
learning_rate	0.0003	
loss	-0.0969	
n_updates	640	
policy_gradient_loss	-0.0641	
value_loss	0.00129	

rollout/		
ep_len_mean	12.6	
ep_rew_mean	0.956	
time/		
fps	718	
iterations	66	
time_elapsed	94	
total_timesteps	67584	
train/		
approx_kl	0.12408055	
clip_fraction	0.23	
clip_range	0.2	
entropy_loss	-0.255	
explained_variance	0.888	
learning_rate	0.0003	
loss	-0.0465	
n_updates	650	
policy_gradient_loss	0.0242	
value_loss	4.94e-05	

rollout/		
ep_len_mean	11.7	
ep_rew_mean	0.959	
time/		
fps	717	

	iterations		67	
	time_elapsed		95	
	total_timesteps		68608	
	train/			
	approx_kl		0.025122331	
	clip_fraction		0.173	
	clip_range		0.2	
	entropy_loss		-0.206	
	explained_variance		0.771	
	learning_rate		0.0003	
	loss		0.0226	
	n_updates		660	
	policy_gradient_loss		-0.0263	
	value_loss		0.000238	

	rollout/			
	ep_len_mean		11.9	
	ep_rew_mean		0.958	
	time/			
	fps		717	
	iterations		68	
	time_elapsed		97	
	total_timesteps		69632	
	train/			
	approx_kl		0.06556613	
	clip_fraction		0.141	
	clip_range		0.2	
	entropy_loss		-0.162	
	explained_variance		0.95	
	learning_rate		0.0003	
	loss		-0.0489	
	n_updates		670	
	policy_gradient_loss		-0.00874	
	value_loss		4.17e-05	

Eval num_timesteps=70000, episode_reward=0.96 +/- 0.00
 Episode length: 11.00 +/- 0.00

	eval/			
	mean_ep_length		11	
	mean_reward		0.961	
	time/			
	total_timesteps		70000	
	train/			
	approx_kl		0.062489584	
	clip_fraction		0.0564	
	clip_range		0.2	
	entropy_loss		-0.0847	
	explained_variance		0.778	
	learning_rate		0.0003	
	loss		-0.0217	
	n_updates		680	
	policy_gradient_loss		-0.0199	
	value_loss		0.000264	

	rollout/			
	ep_len_mean		11.1	
	ep_rew_mean		0.961	
	time/			
	fps		716	
	iterations		69	
	time_elapsed		98	
	total_timesteps		70656	

	rollout/			
	ep_len_mean		13.7	
	ep_rew_mean		0.952	
	time/			
	fps		715	
	iterations		70	
	time_elapsed		100	
	total_timesteps		71680	
	train/			

	approx_kl		0.05116283	
	clip_fraction		0.119	
	clip_range		0.2	
	entropy_loss		-0.131	
	explained_variance		0.985	
	learning_rate		0.0003	
	loss		-0.0654	
	n_updates		690	
	policy_gradient_loss		-0.0128	
	value_loss		9.36e-06	

	rollout/			
	ep_len_mean		13.2	
	ep_rew_mean		0.954	
	time/			
	fps		715	
	iterations		71	
	time_elapsed		101	
	total_timesteps		72704	
	train/			
	approx_kl		0.039104126	
	clip_fraction		0.158	
	clip_range		0.2	
	entropy_loss		-0.24	
	explained_variance		0.178	
	learning_rate		0.0003	
	loss		-0.0193	
	n_updates		700	
	policy_gradient_loss		-0.00839	
	value_loss		0.000591	

	rollout/			
	ep_len_mean		12.7	
	ep_rew_mean		0.956	
	time/			
	fps		715	
	iterations		72	
	time_elapsed		103	
	total_timesteps		73728	
	train/			
	approx_kl		0.058678307	
	clip_fraction		0.166	
	clip_range		0.2	
	entropy_loss		-0.236	
	explained_variance		0.761	
	learning_rate		0.0003	
	loss		0.0961	
	n_updates		710	
	policy_gradient_loss		-0.0254	
	value_loss		0.000299	

	rollout/			
	ep_len_mean		14.2	
	ep_rew_mean		0.95	
	time/			
	fps		714	
	iterations		73	
	time_elapsed		104	
	total_timesteps		74752	
	train/			
	approx_kl		0.042911883	
	clip_fraction		0.147	
	clip_range		0.2	
	entropy_loss		-0.215	
	explained_variance		0.893	
	learning_rate		0.0003	
	loss		-0.0327	
	n_updates		720	
	policy_gradient_loss		-0.0165	
	value_loss		0.000142	

	rollout/			
--	----------	--	--	--

ep_len_mean	14.2	
ep_rew_mean	0.95	
time/		
fps	710	
iterations	74	
time_elapsed	106	
total_timesteps	75776	
train/		
approx_kl	0.02610638	
clip_fraction	0.215	
clip_range	0.2	
entropy_loss	-0.269	
explained_variance	0.556	
learning_rate	0.0003	
loss	-0.0248	
n_updates	730	
policy_gradient_loss	-0.0269	
value_loss	0.000532	

rollout/		
ep_len_mean	13	
ep_rew_mean	0.954	
time/		
fps	710	
iterations	75	
time_elapsed	108	
total_timesteps	76800	
train/		
approx_kl	0.04112082	
clip_fraction	0.249	
clip_range	0.2	
entropy_loss	-0.29	
explained_variance	0.645	
learning_rate	0.0003	
loss	-0.0464	
n_updates	740	
policy_gradient_loss	-0.0367	
value_loss	0.000662	

rollout/		
ep_len_mean	12.1	
ep_rew_mean	0.958	
time/		
fps	711	
iterations	76	
time_elapsed	109	
total_timesteps	77824	
train/		
approx_kl	0.058698382	
clip_fraction	0.252	
clip_range	0.2	
entropy_loss	-0.229	
explained_variance	0.883	
learning_rate	0.0003	
loss	-0.0224	
n_updates	750	
policy_gradient_loss	-0.0408	
value_loss	0.000145	

rollout/		
ep_len_mean	12.5	
ep_rew_mean	0.956	
time/		
fps	711	
iterations	77	
time_elapsed	110	
total_timesteps	78848	
train/		
approx_kl	0.070883654	
clip_fraction	0.226	
clip_range	0.2	
entropy_loss	-0.171	
explained_variance	0.918	

learning_rate	0.0003	
loss	-0.0669	
n_updates	760	
policy_gradient_loss	-0.0423	
value_loss	0.000216	

rollout/		
ep_len_mean	12.2	
ep_rew_mean	0.957	
time/		
fps	712	
iterations	78	
time_elapsed	112	
total_timesteps	79872	
train/		
approx_kl	0.017188694	
clip_fraction	0.0901	
clip_range	0.2	
entropy_loss	-0.171	
explained_variance	0.66	
learning_rate	0.0003	
loss	-0.0169	
n_updates	770	
policy_gradient_loss	-0.0119	
value_loss	0.000467	

Eval num_timesteps=80000, episode_reward=0.96 +/- 0.00
Episode length: 11.00 +/- 0.00

eval/		
mean_ep_length	11	
mean_reward	0.961	
time/		
total_timesteps	80000	
train/		
approx_kl	0.04313474	
clip_fraction	0.107	
clip_range	0.2	
entropy_loss	-0.213	
explained_variance	0.829	
learning_rate	0.0003	
loss	0.00571	
n_updates	780	
policy_gradient_loss	-0.0187	
value_loss	0.000195	

rollout/		
ep_len_mean	12	
ep_rew_mean	0.958	
time/		
fps	712	
iterations	79	
time_elapsed	113	
total_timesteps	80896	

rollout/		
ep_len_mean	12.8	
ep_rew_mean	0.955	
time/		
fps	711	
iterations	80	
time_elapsed	115	
total_timesteps	81920	
train/		
approx_kl	0.06471266	
clip_fraction	0.314	
clip_range	0.2	
entropy_loss	-0.294	
explained_variance	0.938	
learning_rate	0.0003	
loss	-0.0486	
n_updates	790	
policy_gradient_loss	-0.0427	

value_loss	5.16e-05	

rollout/		
ep_len_mean	11.8	
ep_rew_mean	0.958	
time/		
fps	711	
iterations	81	
time_elapsed	116	
total_timesteps	82944	
train/		
approx_kl	0.06405732	
clip_fraction	0.211	
clip_range	0.2	
entropy_loss	-0.176	
explained_variance	0.795	
learning_rate	0.0003	
loss	-0.0243	
n_updates	800	
policy_gradient_loss	-0.0288	
value_loss	0.000221	

rollout/		
ep_len_mean	11.6	
ep_rew_mean	0.959	
time/		
fps	710	
iterations	82	
time_elapsed	118	
total_timesteps	83968	
train/		
approx_kl	0.061312027	
clip_fraction	0.0476	
clip_range	0.2	
entropy_loss	-0.0518	
explained_variance	0.881	
learning_rate	0.0003	
loss	-0.0393	
n_updates	810	
policy_gradient_loss	-0.023	
value_loss	0.000125	

rollout/		
ep_len_mean	12.2	
ep_rew_mean	0.957	
time/		
fps	708	
iterations	83	
time_elapsed	119	
total_timesteps	84992	
train/		
approx_kl	0.039377097	
clip_fraction	0.181	
clip_range	0.2	
entropy_loss	-0.117	
explained_variance	0.814	
learning_rate	0.0003	
loss	-0.0468	
n_updates	820	
policy_gradient_loss	0.022	
value_loss	0.000187	

rollout/		
ep_len_mean	11.8	
ep_rew_mean	0.958	
time/		
fps	708	
iterations	84	
time_elapsed	121	
total_timesteps	86016	
train/		
approx_kl	0.022505693	

clip_fraction	0.109	
clip_range	0.2	
entropy_loss	-0.158	
explained_variance	0.671	
learning_rate	0.0003	
loss	-0.0147	
n_updates	830	
policy_gradient_loss	-0.0127	
value_loss	0.000268	

rollout/		
ep_len_mean	15.4	
ep_rew_mean	0.946	
time/		
fps	707	
iterations	85	
time_elapsed	123	
total_timesteps	87040	
train/		
approx_kl	0.11171889	
clip_fraction	0.142	
clip_range	0.2	
entropy_loss	-0.131	
explained_variance	0.914	
learning_rate	0.0003	
loss	-0.0735	
n_updates	840	
policy_gradient_loss	-0.0356	
value_loss	6.9e-05	

rollout/		
ep_len_mean	14.4	
ep_rew_mean	0.95	
time/		
fps	706	
iterations	86	
time_elapsed	124	
total_timesteps	88064	
train/		
approx_kl	0.05634898	
clip_fraction	0.322	
clip_range	0.2	
entropy_loss	-0.321	
explained_variance	0.337	
learning_rate	0.0003	
loss	-0.0435	
n_updates	850	
policy_gradient_loss	-0.03	
value_loss	0.00163	

rollout/		
ep_len_mean	11.6	
ep_rew_mean	0.959	
time/		
fps	706	
iterations	87	
time_elapsed	126	
total_timesteps	89088	
train/		
approx_kl	0.1144023	
clip_fraction	0.191	
clip_range	0.2	
entropy_loss	-0.188	
explained_variance	0.632	
learning_rate	0.0003	
loss	-0.0365	
n_updates	860	
policy_gradient_loss	-0.0271	
value_loss	0.000624	

Eval num_timesteps=90000, episode_reward=0.96 +/- 0.00
Episode length: 11.00 +/- 0.00

eval/		
mean_ep_length	11	
mean_reward	0.961	
time/		
total_timesteps	90000	
train/		
approx_kl	0.065284915	
clip_fraction	0.063	
clip_range	0.2	
entropy_loss	-0.0858	
explained_variance	0.89	
learning_rate	0.0003	
loss	-0.0338	
n_updates	870	
policy_gradient_loss	0.0276	
value_loss	0.000171	

rollout/		
ep_len_mean	12.8	
ep_rew_mean	0.955	
time/		
fps	705	
iterations	88	
time_elapsed	127	
total_timesteps	90112	

rollout/		
ep_len_mean	14.6	
ep_rew_mean	0.949	
time/		
fps	704	
iterations	89	
time_elapsed	129	
total_timesteps	91136	
train/		
approx_kl	0.096123055	
clip_fraction	0.196	
clip_range	0.2	
entropy_loss	-0.191	
explained_variance	0.633	
learning_rate	0.0003	
loss	-0.0416	
n_updates	880	
policy_gradient_loss	-0.0249	
value_loss	0.000407	

rollout/		
ep_len_mean	11.1	
ep_rew_mean	0.961	
time/		
fps	704	
iterations	90	
time_elapsed	130	
total_timesteps	92160	
train/		
approx_kl	0.19162205	
clip_fraction	0.145	
clip_range	0.2	
entropy_loss	-0.15	
explained_variance	0.256	
learning_rate	0.0003	
loss	0.0508	
n_updates	890	
policy_gradient_loss	-0.0202	
value_loss	0.00348	

rollout/		
ep_len_mean	12.3	
ep_rew_mean	0.957	
time/		
fps	703	
iterations	91	

time_elapsed	132	
total_timesteps	93184	
train/		
approx_kl	0.06836694	
clip_fraction	0.347	
clip_range	0.2	
entropy_loss	-0.264	
explained_variance	0.581	
learning_rate	0.0003	
loss	-0.0699	
n_updates	900	
policy_gradient_loss	0.066	
value_loss	0.000154	

rollout/		
ep_len_mean	12.4	
ep_rew_mean	0.957	
time/		
fps	703	
iterations	92	
time_elapsed	133	
total_timesteps	94208	
train/		
approx_kl	0.020714343	
clip_fraction	0.161	
clip_range	0.2	
entropy_loss	-0.218	
explained_variance	0.581	
learning_rate	0.0003	
loss	-0.01	
n_updates	910	
policy_gradient_loss	-0.00672	
value_loss	0.000254	

rollout/		
ep_len_mean	12.3	
ep_rew_mean	0.957	
time/		
fps	703	
iterations	93	
time_elapsed	135	
total_timesteps	95232	
train/		
approx_kl	0.037652392	
clip_fraction	0.184	
clip_range	0.2	
entropy_loss	-0.199	
explained_variance	0.896	
learning_rate	0.0003	
loss	-0.0423	
n_updates	920	
policy_gradient_loss	-0.0175	
value_loss	0.000112	

rollout/		
ep_len_mean	11.2	
ep_rew_mean	0.961	
time/		
fps	702	
iterations	94	
time_elapsed	136	
total_timesteps	96256	
train/		
approx_kl	0.080989845	
clip_fraction	0.205	
clip_range	0.2	
entropy_loss	-0.133	
explained_variance	0.918	
learning_rate	0.0003	
loss	-0.05	
n_updates	930	
policy_gradient_loss	-0.0212	
value_loss	8.25e-05	

```

-----
| rollout/          |          |
| ep_len_mean      | 21       |
| ep_rew_mean      | 0.923    |
| time/            |          |
| fps              | 703      |
| iterations        | 95       |
| time_elapsed      | 138      |
| total_timesteps   | 97280    |
| train/           |          |
| approx_kl         | 0.13105541 |
| clip_fraction     | 0.316    |
| clip_range        | 0.2       |
| entropy_loss      | -0.254    |
| explained_variance | 0.971     |
| learning_rate      | 0.0003    |
| loss              | -0.0671   |
| n_updates         | 940       |
| policy_gradient_loss | 0.197     |
| value_loss        | 1.27e-05  |
-----

```

```

-----
| rollout/          |          |
| ep_len_mean      | 27.9     |
| ep_rew_mean      | 0.897     |
| time/            |          |
| fps              | 703      |
| iterations        | 96        |
| time_elapsed      | 139      |
| total_timesteps   | 98304    |
| train/           |          |
| approx_kl         | 0.1051268 |
| clip_fraction     | 0.194     |
| clip_range        | 0.2       |
| entropy_loss      | -0.144    |
| explained_variance | -0.119    |
| learning_rate      | 0.0003    |
| loss              | -0.0391   |
| n_updates         | 950       |
| policy_gradient_loss | 0.162     |
| value_loss        | 0.00478   |
-----

```

```

-----
| rollout/          |          |
| ep_len_mean      | 15.9     |
| ep_rew_mean      | 0.943     |
| time/            |          |
| fps              | 704       |
-----

```

```

-----
| iterations        | 97        |
| time_elapsed      | 140       |
| total_timesteps   | 99328     |
| train/           |          |
| approx_kl         | 0.5337625 |
| clip_fraction     | 0.366     |
| clip_range        | 0.2        |
| explained_variance | -0.0747   |
| learning_rate      | 0.0003    |
| learning_rate      | 0.0003    |
| loss              | -0.0265   |
| n_updates         | 960       |
| policy_gradient_loss | 0.308     |
| value_loss        | 0.00451   |
-----

```

Eval num_timesteps=100000, episode_reward=0.96 +/- 0.00
 Episode length: 11.00 +/- 0.00

```

-----
| eval/            |          |
| mean_ep_length    | 11        |
| mean_reward       | 0.961     |
| time/            |          |
| total_timesteps   | 100000    |
| train/           |          |
| approx_kl         | 0.13666381 |
| clip_fraction     | 0.168     |
| clip_range        | 0.2        |
| entropy_loss      | -0.155     |
| explained_variance | -0.427     |
| learning_rate      | 0.0003    |
| loss              | -0.0696   |
| n_updates         | 970       |
| policy_gradient_loss | 0.0638     |
| value_loss        | 0.00239   |
-----

```

```

-----
| rollout/          |          |
| ep_len_mean      | 13.1      |
| ep_rew_mean      | 0.954     |
| time/            |          |
| fps              | 705       |
| iterations        | 98        |
| time_elapsed      | 142       |
| total_timesteps   | 100352    |
-----

```

Training complete.
 Model saved as ppo_minigrid_model.
 Testing trained agent...

APPENDIX D

Model Trained in SageMaker Instance

The images are captured during the training process in AWS SageMaker and logs are stored in the S3 bucket.

The screenshot shows a JupyterLab interface within a web browser. The address bar displays the URL: `https://projectcheck.notebook.us-east-1.sagemaker.aws/tree/logs`. The JupyterLab header includes the 'jupyter' logo and buttons for 'Open JupyterLab', 'Quit', and 'Logout'. Below the header, there are tabs for 'Files', 'Running', 'Clusters', 'SageMaker Examples', and 'Conda'. The 'Files' tab is active, showing a file explorer for the '/logs' directory. It lists several files: '..', 'best_model', 'evaluations.npz', 'events.out.tfevents.1732859462.ip-172-16-117-137.ec2.internal.17519.0', and 'progress.csv', along with their last modified times and file sizes. Below the file explorer, there are two code input areas. The first, labeled 'In [7]:', contains a Python function `def test_agent(model, env):` that tests a trained RL agent. The second, labeled 'In [8]:', contains code to create an environment and train an agent. Below the code, the output console displays training metrics in a table format, including entropy_loss, explained_variance, learning_rate, loss, n_updates, policy_gradient_loss, value_loss, rollout/ep_len_mean, rollout/ep_rew_mean, time/fps, iterations, time_elapsed, and total_timesteps. At the bottom of the output, a message states: 'Model saved to S3: s3://awsaisoccertest/models/ppo_minigrid_2024-11-29_05-51-01.zip'.

```
In [7]: def test_agent(model, env):
        """Tests a trained RL agent on the given environment."""
        obs, _ = env.reset()
        done = False
        while not done:
            env.render()
            action, _ = model.predict(obs)
            obs, reward, done, info = env.step(action)
            print(f"Reward: {reward}")

In [8]: env = create_environment()
        model = train_agent(env)
```


entropy_loss	-0.134
explained_variance	0.615
learning_rate	0.0003
loss	-0.0335
n_updates	970
policy_gradient_loss	-0.0233
value_loss	0.000439

rollout/	
ep_len_mean	12.6
ep_rew_mean	0.956
time/	
fps	364
iterations	98
time_elapsed	275
total_timesteps	100352

Model saved to S3: s3://awsaisoccertest/models/ppo_minigrid_2024-11-29_05-51-01.zip

```
Episode finished after 7 steps with total reward: 0.9369999999999999
Episode finished after 34 steps with total reward: 0.694
Episode finished after 100 steps with total reward: 0
Episode finished after 26 steps with total reward: 0.766
Episode finished after 100 steps with total reward: 0
Episode finished after 100 steps with total reward: 0
Episode finished after 100 steps with total reward: 0
Episode finished after 27 steps with total reward: 0.757
Episode finished after 100 steps with total reward: 0
Episode finished after 100 steps with total reward: 0
```

Test Results



```
https://projectgpu.notebook.us-east-1.sagemaker.aws/edit/logs/progress.csv
jupyter progress.csv 23 minutes ago
File Edit View Language current mode

15 3,13312,,13,3567,,0.01550966,120,-0.0074328081944258885,9.810567634715993e-
07,0.2,0.09248046875,-1.59291105940938,0.0003,-0.04184097424149513,-1.8361585140228271,0.031232,2468.8,,
16 3,14336,,14,3709,,0.01339952,130,-0.002147884003352374,4.243812977478001e-
07,0.2,0.03095703125,-1.5994115896522998,0.0003,-0.03436832129955292,-0.8418478965759277,0.031232,2468.8,,
17 3,15360,,15,3852,,0.012507655,140,-0.005490707751596347,6.006332903041312e-
07,0.2,0.07412109375,-1.6025339804589749,0.0003,-0.015489249490201473,-1.3051090240478516,0.02602666666666667,2474.0,,
18 4,16384,,16,3992,,0.011104648,150,-0.004503814497729764,3.3050289109937126e-
07,0.2,0.05244140625,-1.58695924654603,0.0003,-0.026387197896838188,-3.6042418479919434,0.02602666666666667,2474.0,,
19 4,17408,,17,4132,,0.012377972,160,-0.007170661754207686,2.625206301676286e-
07,0.2,0.069921875,-1.5983815297484398,0.0003,-0.013051476329565048,0.03473716974258423,0.022308571428571428,2477.714285714286,,
20 4,18432,,18,4271,,0.013479005,170,-0.003753540207981132,1.4068342411022172e-
07,0.2,0.03798828125,-1.5966101825237273,0.0003,-0.0016077914042398334,-1.1010208129882812,0.022308571428571428,2477.714285714286,,
21 4,19456,,19,4419,,0.011153008,180,-0.0046226314065279436,2.2237360479682168e-
07,0.2,0.05576171875,-1.5802074290812016,0.0003,-0.02338787168264389,-0.01065516471862793,0.022308571428571428,2477.714285714286,,
22 3,20000,,20,4559,,0.01579111,190,-0.007130152359604835,2.8066454481331715e-
07,0.2,0.07509765625,-1.548928889632225,0.0003,-0.04514220729470253,-2.3442389965057373,,0.0,2500.0
23 3,20480,,20,6257,,0.012394222222222223,2217.4444444444443,,
24 3,21504,,21,6398,,0.0146340355,200,-0.005744844989385456,0.00812491550918253,0.2,0.14921875,-1.4724476367235184,0.0003,-0.0199076496064662
93,-0.001684069633488867,0.12394222222222223,2217.4444444444443,,
25 3,22528,,22,6543,,0.00957765,210,-0.0036235730891348793,7.903866219294287e-
05,0.2,0.03740234375,-1.4574183650314807,0.0003,-0.0014033853076398373,-3.704458236694336,0.143292,2185.3,,
26 3,23552,,23,6684,,0.013656424,220,-0.004584695975063368,0.001038198106854793,0.2,0.11455078125,-1.4145512960851192,0.0003,0.00530810607597
2319,-0.001232147216796875,0.143292,2185.3,,
27 3,24576,,24,6828,,0.009005114,230,-0.003219434808124788,1.9320533059877244e-
08,0.2,0.04599609375,-1.4208344288170338,0.0003,0.01128444168716669,-1.6931993961334229,0.13026545454545455,2213.909090909091,,
28 3,25600,,25,6968,,0.009980627,240,-0.006319787554093637,2.3737630840126035e-
09,0.2,0.083203125,-1.3963280409574508,0.0003,-0.010413549840450287,-0.9580490589141846,0.13026545454545455,2213.909090909091,,
29 3,26624,,26,7109,,0.009770433,250,-0.005123643798287958,1.4027877730943272e-
09,0.2,0.0982421875,-1.44269190877676,0.0003,-0.02505362033843994,-3.3265914916992188,0.13026545454545455,2213.909090909091,,
30 3,27648,,27,7252,,0.012489204,260,-0.0017364906822940409,1.269701439127101e-
09,0.2,0.03798828125,-1.4804138235747815,0.0003,-0.011584128253161907,-2.849370002746582,0.11941,2237.75,,
31 3,28672,,28,7392,,0.0074411863,270,-0.006641380654764361,2.829135061638288e-
09,0.2,0.0728515625,-1.4820755116641522,0.0003,-0.015028017573058605,-1.825103521347046,0.11941,2237.75,,
32 3,29696,,29,7535,,0.015512468,280,-0.0035945162846473975,1.945664283853433e-
09,0.2,0.051953125,-1.4549897141754626,0.0003,0.006157361902296543,-6.177978992462158,0.11022461538461538,2257.923076923077,,
33
```

General purpose buckets

Directory buckets

General purpose buckets (3) Info All AWS Regions

Copy ARN

Empty

Delete

Create bucket

Buckets are containers for data stored in S3.

Find buckets by name

< 1 >

Name	AWS Region	IAM Access Analyzer	Creation date
awsaisoccertest	US East (N. Virginia) us-east-1	View analyzer for us-east-1	November 28, 2024, 21:07:32 (UTC-08:00)

Amazon S3

Buckets

Access Grants

Access Points

Object Lambda Access Points

Multi-Region Access Points

Batch Operations

IAM Access Analyzer for S3

Block Public Access settings for this account

Storage Lens

Dashboards

Storage Lens groups

awsaisoccertest info

Objects

Properties

Permissions

Metrics

Management

Access Points

Objects (1) Info

Copy S3 URI

Copy URL

Download

Open

Delete

Actions

Create folder

Upload

Find objects by prefix

Name	Type	Last modified	Size	Storage class
models/	Folder	-	-	-

models/

Copy S3 URI

Objects

Properties

Objects (1) Info

Copy S3 URI

Copy URL

Download

Open

Delete

Actions

Create folder

Upload

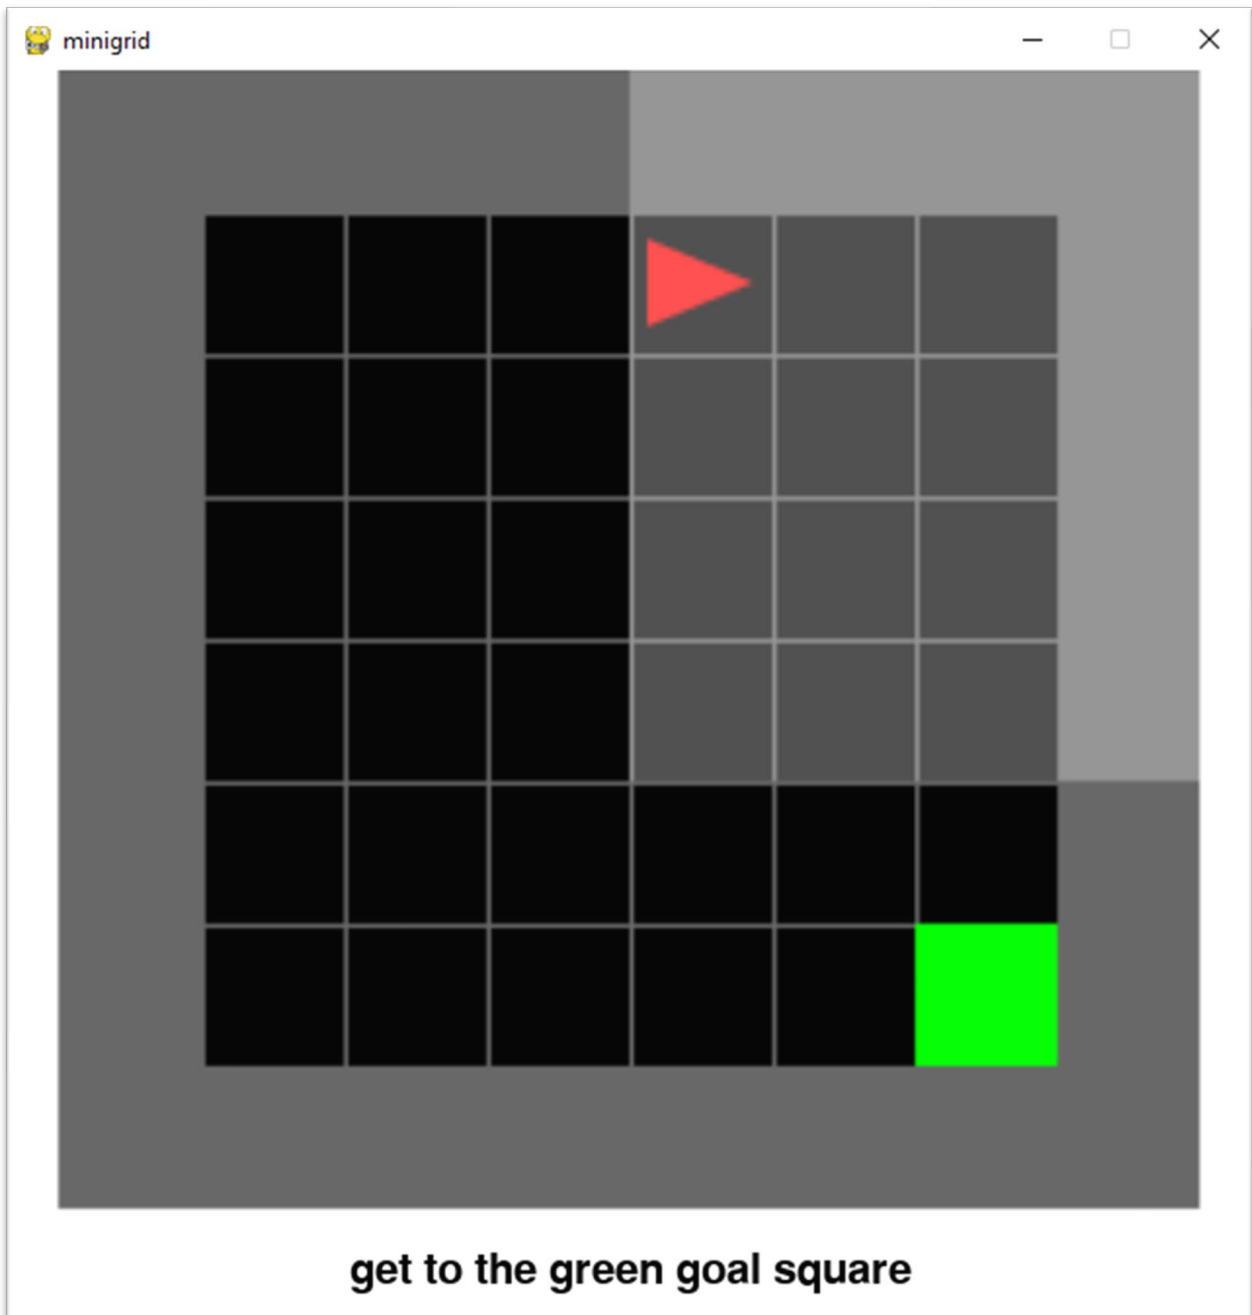
Find objects by prefix

Name	Type	Last modified	Size	Storage class
ppo_minigrid_2024-11-29_05-51-01.zip	zip	November 28, 2024, 21:55:43 (UTC-08:00)	432.9 KB	Standard

APPENDIX E

Render Mode = Human

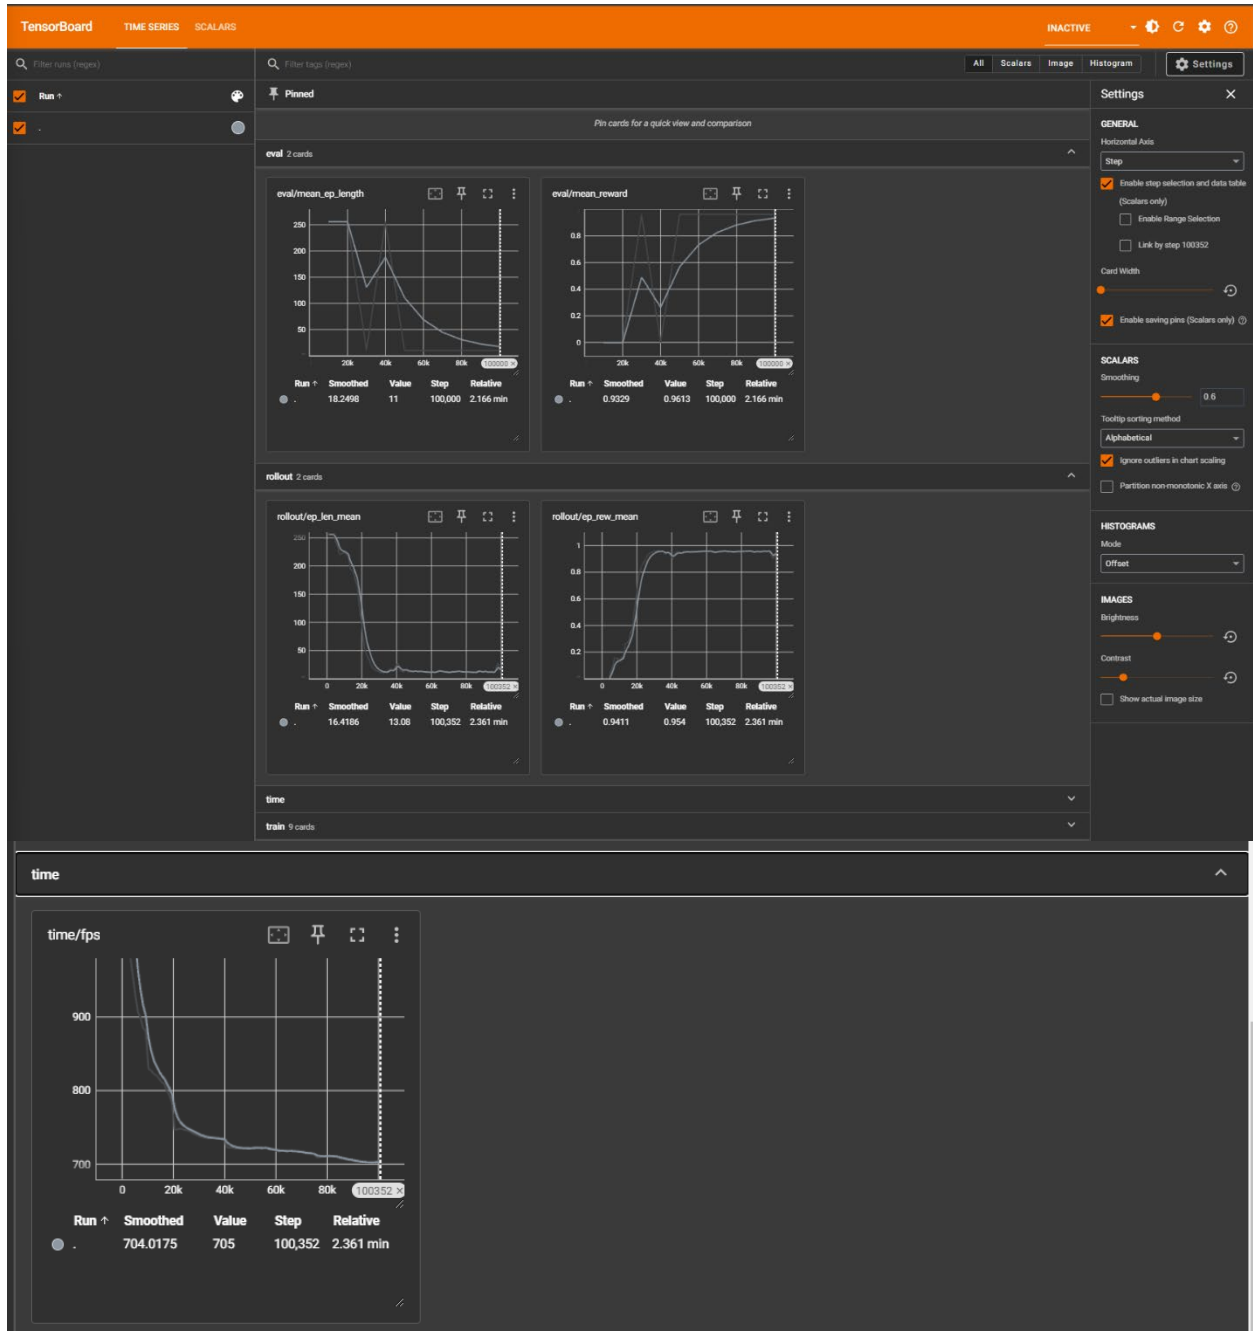
This still image, captured during training in human render mode, illustrates the agent (red triangle) navigating a grid environment toward the green goal square as per the specified task.

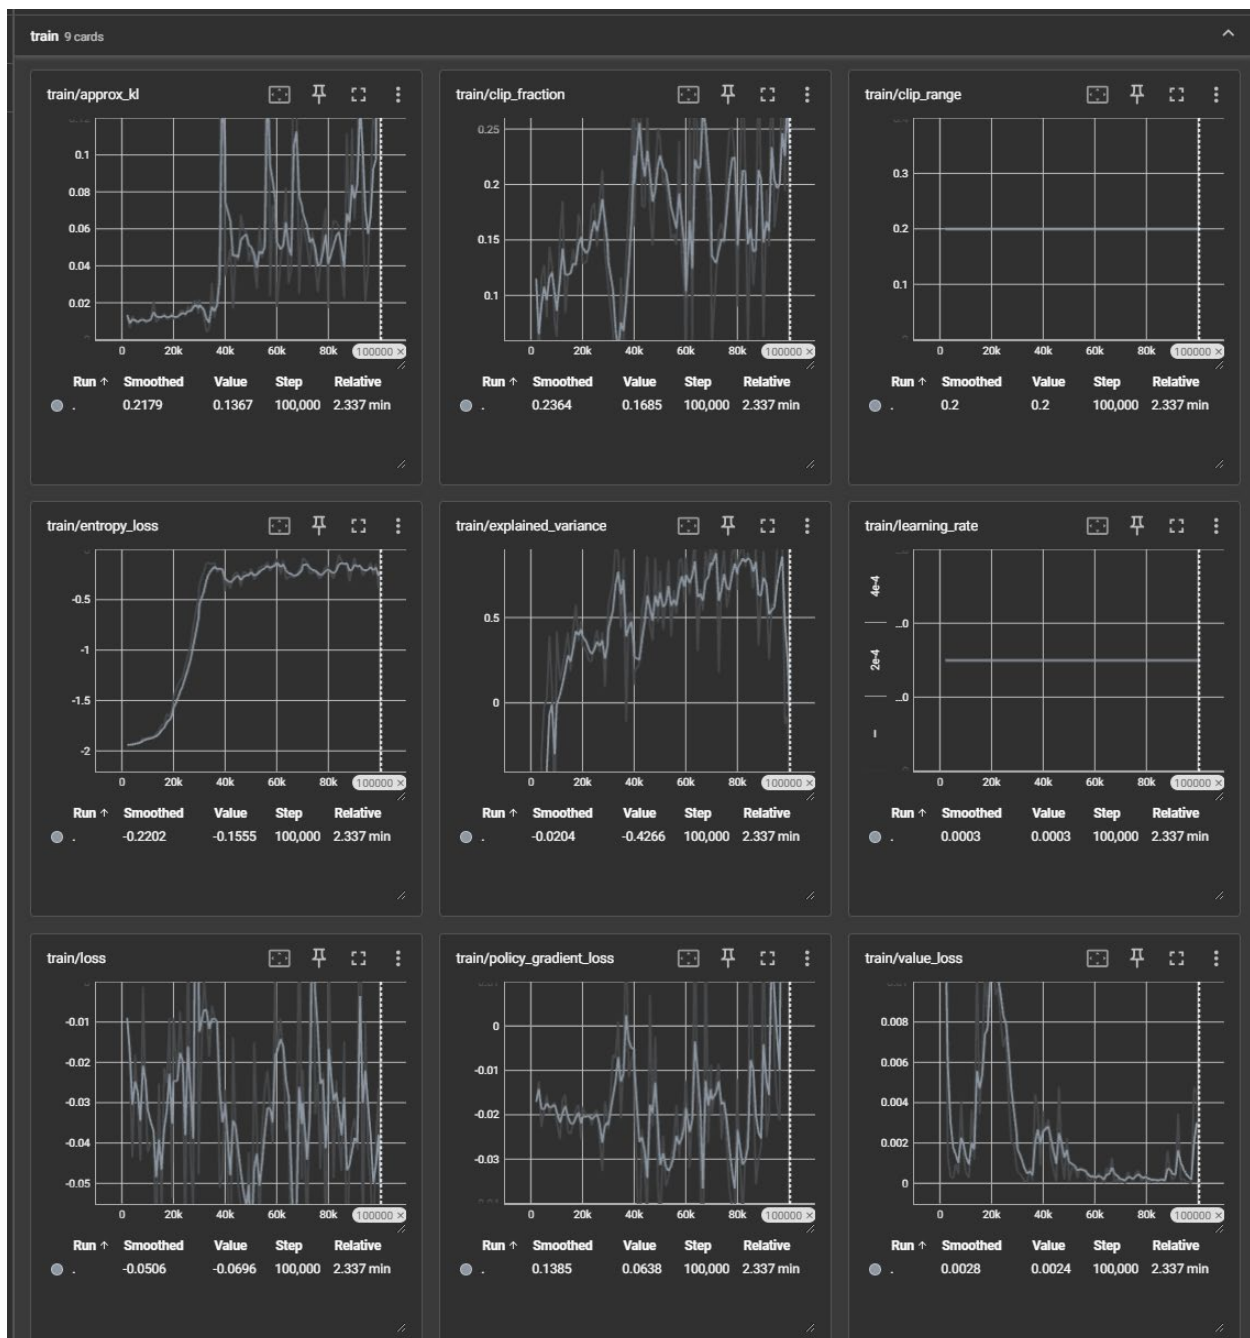


APPENDIX F

TensorBoard Logs – 1st Visualized Training Metrics

This appendix contains detailed logs from the initial evaluation phase, highlighting the agent's performance metrics, episode outcomes, and termination conditions during validation runs.





APPENDIX G

TensorBoard Logs – 2nd Visualized Training Metrics + Additional Details

This appendix provides comprehensive logs from the training runs, documenting the agent's progress, action sequences, rewards, and termination conditions to support the analysis of training performance.

